

Dynamic sampling pointnet notes

xyz

Feb 2018

Contents

1 Quick notes for important events while using one file to test	1
1.1 batch size	1
1.1.1 bs=27 vs bs=81	1
1.2 feed elements	2
1.3 model	3
1.4 integration: matterport3d	4
1.5 integration: scannet	6
1.6 Semantic segmentation expamples	6
1.6.1 good: 1083, train, 0.946	6
1.6.2 bad: 18737,eval 0.071	6
1.7 point++	7
1.7.1 scannet seg	7

1 Quick notes for important events while using one file to test

1.1 batch size

1.1.1 bs=27 vs bs=81

batch size: 9,27,81

data: xyz-color_1norm

model: 1AG

sampling & grouping: stride_0d1_step_0d1_bmap_nh5_2048_0d5_1_fmnl-160_32-32_12-0d2_0d6-0d2_0d6

Figure 1: bs=9





Figure 2: bs=27



Figure 3: bs=81



1.2 feed elements

epoch num = 100

stride_0d1_step_0d1_bmap_nh5.2048.0d5.1_fmnl-160_32-32_12-0d2_0d6-0d2_0d6

model	batch size	data elements	acc	loss
1AG	9	xyz color	0.890	0.356
1AG	27	xyz color	0.920	0.240
3AG	27	xyz color	0.912	0.273
2A	27	xyz color	0.908	0.294
2AG	27	xyz color	0.902	0.293
1A	27	xyz color	0.883	0.351
1AG	81	xyz color	0.978	0.072
1AG	9	xyz	0.861	0.427
1AG	27	xyz	0.907	0.257
1AG	81	xyz	0.975	0.078
1A	27	xyzmid color	0.889	0.357
3AG	27	xyzmid color	0.933	0.193
2A	27	xyzmid color	0.939	0.177
2AG	27	xyzmid color	0.929	0.208
3AG	27	xyz xyzmid color	0.924	0.230
2A	27	xyz xyzmid color	0.898	0.317
2AG	27	xyz xyzmid color	0.908	0.280
1A	27	xyz xyzmid color	0.910	0.281
1AG	27	xyz xyzmid color	0.944	0.163
1AG	81	xyz xyzmid color	0.976	0.078
2A	81	xyz xyzmid color	0.942	0.173
3AG	81	xyz xyzmid color	0.949	0.147

1. large batch size is better
2. $1AG(0.92) > 3AG(0.912) > 2A(0.908) > 2AG(0.902) > 1A(883)$
1AG is much better than 1A
1AG is a bit better than 3AG ???
3. xyz-color is only a bit better than xyz
4. xyzmid-color is much better than xyz-color
5. **xyzmid-color is normally much better than xyz-xyzmid-color**
???

1.3 model

batch size: 50

data: xyz_midnorm_block-color_1norm

epoch_num = 600

sampling & grouping: stride_0d1_step_0d1_bmap_nh5_12800_1d6_2_fmn3-600_64_24-60_16_12-0d2_0d6_1d2-0d2_0d6_1d2

model	acc	loss
3A	0.909	0.248
3AG	0.913	0.231
4AG	0.912	0.232

batch size: 32
 data: xyz_midnorm_block-color_1norm
 sampling & grouping: stride_0d1_step_0d1_bmap_nh5_12800_1d6_2_fmn6-2048_256_64-32_32_16-0d2_0d6_1d2-0d1_0d3_0d6
 matterport3d
 feed_data_elements:['xyz_midnorm_block', 'color_1norm']
 feed_label_elements:['label_category', 'label_instance']
 train data shape: [362 12800 6]
 test data shape: [384 12800 6]
 max epoch = 500

model	acc	loss
1AG	0.944/0.431	0.161/4.633
4AG	0.835/0.401	0.520/3.644

1.4 integration: matterport3d

stride_0d1_step_0d1_bmap_nh5_12800_1d6_2_fm3-512.64.24-48.16.12-0d2_0d6_1d2-0d2_0d6_1d2				
17D_1LX_1pX_29h_2az				
model	batch size batch num shuffle	lr ds	data elements	epoch-acc mean-std train/eval
1aG	30/60	0.005	'xyz_midnorm_block', 'color_1norm', 'nxnynz'	250-0.981
1DSaG	30/60	0.001-40	'xyz_midnorm_block', 'color_1norm', 'nxnynz'	300-0.914-0.775
1DSaG	30/60	0.001-40	'xyz_midnorm_block', 'color_1norm', 'nxnynz'	300-0.914-0.775
1DSaG kp0.5	30/60	0.001-80 300-3e-4	'xyz_midnorm_block', 'color_1norm', 'nxnynz'	300-0.942-0.842
1DSaG kp0.2	30/60	0.001-80 300-3e-4	'xyz_midnorm_block', 'color_1norm', 'nxnynz'	300-0.928-0.797
1DSaG kp0.5	30/60	0.005-80 300-1.7e-3	'xyz_midnorm_block', 'color_1norm', 'nxnynz'	300-0.970-0.916
1DSaG kp0.2	30/60	0.005-80 300-1.7e-3	'xyz_midnorm_block', 'color_1norm', 'nxnynz'	300-0.966-0.924
1DSaG kp0.8	30/60	0.005-80 300-1.7e-3	'xyz_midnorm_block', 'color_1norm', 'nxnynz'	300-0.976-0.933 500-0.984-0.954
1aG	30/1083	0.003	'xyz_midnorm_block', 'color_1norm', 'nxnynz'	200-0.947
1aG	30/1083	0.01	'xyz_midnorm_block', 'color_1norm'	200-0.783 500-0.791
1aG	30/1083	0.003/30 300-0.00012	'xyz_midnorm_block', 'color_1norm'	200-0.903 300-0.921
1bG	25/1083	0.001-30 100-3e-4 300-4e-5	'xyz_midnorm_block'	100-0.854 200-0.918 300-0.936
1bG	25/1083	0.001-30 100-3e-4 300-4e-5	'xyz_midnorm_block', 'color_1norm', 'nxnynz'	100-0.914 200-0.957 300-0.966
1bG	25/1083	0.02	'xyz_midnorm_block', 'color_1norm'	200-0.655 300-0.718
1bG	25/1083	0.02	'xyz_midnorm_block', 'color_1norm', 'nxnynz'	200-0.772 300-0.823
1bG	25/1083	0.001	'xyz'	200-0.772 90-0.553-0.210
4bG	25/1083	0.001-30 100-3e-4 200-1e-4 300-4e-5	'xyz_midnorm_block', 'color_1norm', 'nxnynz'	100-0.752 200-0.816 300-0.832
1aG	30/19755	0.001-30 50-7e-4 100-3e-4	'xyz_midnorm_block', 'color_1norm', 'nxnynz'	50-0.752/0.580 100-0.843/0.574 (NoShuf) 102-0.806/0.570 (Shuffle)
1bG	25/19755	0.001-30	'xyz_midnorm_block'	38-0.719/0.587

1.5 integration: scannet

stride_0d1_step_0d1_bmap_nh5_12800_1d6_2_fm3-256_48_16-56_8-8-0d2_0d6_1d2-0d2_0d6_1d2					
scannet train					
model	loss: E,N,C	batch size batch num shuffle	lr ds	data elements	epoch-point ac-class ac train/eval
1bG	E	25/12887 Y	0.001 40	xyzmid	23-0.732-0.326/0.664-0.260 25-0.746-0.340/0.669-0.273
1bG	N	25/12887 Y	0.001 40	xyzmid	25-0.733-0.390/0.666-0.252
1bG	C	25/12887 Y	0.001 40	xyzmid	25-0.703-0.356/0.655-0.252
1bG	CN	25/12887 Y	0.001 40	xyzmid	25-0.681-0.366/0.611-0.237
4bG	CN	25/2998- 3521 Y	0.001 40	xyzmid	142-0.726-0.445/0.625-0.242
4bG	E	25/2998- 3521 Y	0.001 40	xyzmid	145-0.792-0.506/0.656-0.257

1.6 Semantic segmentation examples

1.6.1 good: 1083, train, 0.946

log: log-model_1bG-gsbb_3B1-bs25-lr1-ds_30-xyz_midnorm_block-color_1norm-nxnynz-12800-mat_1083

model: 1bG

sampling & grouping:

stride_0d1_step_0d1_bmap_nh5_12800_1d6_2_fm3-512_64_24-48_16_12-0d2_0d6_1d2-0d2_0d6_1d2

batch size: 25

learning rate: 0.001000

decay_epoch_step: 30

matterport3d

feed_data_elements:['xyz_midnorm_block', 'color_1norm', 'nxnynz']

feed_label_elements:['label_category', 'label_instance']

train data shape: [1083 12800 9]

1.6.2 bad: 18737,eval 0.071


model: 1bG

sampling & grouping: stride_0d1_step_0d1_bmap_nh5_12800_1d6_2_fmn3-512_64_24-48_16_12-0d2_0d6_1d2-0d2_0d6_1d2
 batch size: 25
 learning rate: 0.001000
 decay_epoch_step: 50
 epoch 0 train IsShuffleIdx: True
 epoch 0 train IsShuffleIdx: True
 matterport3d
 feed_data_elements:['xyz_midnorm_block', 'color_1norm', 'nxnynz']
 feed_label_elements:['label_category', 'label_instance']
 train data shape: [18737 12800 9]
 test data shape: [4172 12800 9]


1.7 point++

1.7.1 scannet seg


each room as a block, total 40 block			
batch size batch num	lr ds	data elements	epoch-point ac-class ac train/eval/eval whole scene
30/40	0.001	xyzmid	200-0.675/0.757-0.54/0.799-0.52
25	0.001	xyzmid	200-0.689/0.787-0.556/0.815-0.517i




(a) colorized point cloud




(b) raw point cloud



(c) gt




(d) pred




(e) err

Figure 4: 17DRP5sb8fy_1_2_a946




(a) colorized point cloud



(b) raw point cloud



(c) gt



(d) pred



(e) err

Figure 5: 17DRP5sb8fy_0_25_a946



Figure 6: qoi_r1Q_r47_rPc_rqf_2_3_a0d071 (raw,gt,pred,err,crt)